American Economic Review: Papers & Proceedings 2017, 107(5): 322–326
https://doi.org/10.1257/aer.p20171020

ENTREPRENEURSHIP, FIRM DYNAMICS, AND GROWTH

Declining Dynamism, Allocative Efficiency, and the Productivity Slowdown†

By Ryan A. Decker, John Haltiwanger, Ron S. Jarmin, and Javier Miranda*

Evidence of declining entrepreneurship and labor market fluidity has captured wide interest among researchers and policymakers. Startup rates and other measures of young firm activity have declined since the 1980s, with accelerated slowdowns in high-growth young firm activity since 2000. Gross job and worker flows have declined over the same period including marked drops since the early 2000s. These patterns are particularly notable in the high-tech sector, which saw rising dynamism during the 1990s before declining sharply after 2000 (Decker et al. 2016).

A distinct literature describes a decline in the growth rate of aggregate productivity since the early 2000s (Byrne, Fernald, and Reinsdorf 2016; Gordon 2016; Syverson 2016). An important omission from much of the productivity slowdown literature is the notion that aggregate productivity depends not only on technology but also on allocative efficiency—the continual movement of resources to their most productive uses. Decker et al. (2017) document declining establishment-level responsiveness of growth to productivity and show that a weakening of the growth-productivity relationship at the business level has had potentially large implications for aggregate productivity growth while also helping explain falling rates of job reallocation. Moreover, within-industry dispersion of labor productivity—a popular (if limited) indicator of efficiency frictions—has risen since the late 1990s. In the present study, we provide further evidence linking the problem of slowing productivity growth to declining business dynamism.

Using firm-level data on labor productivity, we construct accounting decompositions to show that dampened growth in allocative efficiency can account for much of the decline in aggregate productivity growth between the late 1990s and the mid-2000s; more specifically, the slowdown reflects inefficient allocation of productive resources as well as the interaction between allocation and slowing within-firm productivity growth. Our findings imply that, consistent with the conclusions of Decker et al. (2017), declining business dynamism since 2000 is likely a drag on American living standards. Moreover, our findings suggest a reevaluation of the productivity slowdown debate, which has until now focused on technological versus measurement explanations.

I. A Microdata Approach

Our dataset, the RE-LBD, combines the industry and employment data of the Census Bureau's Longitudinal Business Database (LBD) with revenue data from tax records (Haltiwanger et al. forthcoming). The integrated data allow us to measure gross revenue labor productivity at the firm level for virtually

* Decker: Board of Governors of the Federal Reserve, Mailstop 82, 20th Street and C Street NW, Washington, DC 20551 (e-mail: ryan.a.decker@frb.gov); Haltiwanger: Department of Economics, University of Maryland, College Park, MD 20742 (e-mail: haltiwan@econ.umd.edu); Jarmin: US Census Bureau, 4600, Silver Hill Road, Washington, DC 20233 (e-mail: ron.s.jarmin@census.gov); Miranda: US Census Bureau, 4600 Silver Hill Road, Washington, DC 20233 (e-mail: javier.miranda@census.gov). We gratefully acknowledge financial support from the Kauffman Foundation. Cody Tuttle provided excellent research assistance. Any opinions and conclusions expressed herein are those of the authors and do not necessarily represent the views of the US Census Bureau or of the Board of Governors or its staff. All results have been reviewed to ensure that no confidential information is disclosed.

† Go to https://doi.org/10.1257/aer.p20171020 to visit the article page for additional materials and author disclosure statement(s).

the entire US private nonfarm economy; we apply propensity score weights to account for imperfect match rates between revenue and employment (LBD) data.

We construct aggregate labor productivity numbers that come reasonably close to official figures published by the BLS, which rely on different methodology and source data. BLS numbers are based on value added per worker, while we are limited to gross revenue per worker (deflated with BEA deflators, typically at the three-digit or four-digit NAICS level). Earlier research (e.g., Foster, Haltiwanger, and Krizan 2001) shows that gross output per worker tracks value added per worker reasonably well within industries but poorly across industries. We therefore focus on variation within detailed (six-digit NAICS) industries. We construct industry-level indices by aggregating the firm-level data using the employment-weighted average of log labor productivity: $P_{it} = \sum_{f \in i} \theta_{ft} p_{ft}$, where P_{it} is industry-level productivity for industry i in year t, θ_{ft} is the share of employment for firm f in year t, and p_{ft} is log labor productivity for firm f in year t . We aggregate our industry-level computations to an economywide level using fixed industry weights (reflecting each industry's average weight for the entire time sample), thereby avoiding inferences based on cross-industry variation in gross output per worker.

Despite differences from BLS methodology and source data, we obtain similar patterns of aggregate labor productivity growth. In Figure A1 of the online Appendix we report average annual log differences of aggregate productivity from both the RE-LBD and BLS data for three periods: 1997–1999, 2004–2006, and 2011–2013. We report business cycle peaks to avoid cyclical issues; we will focus on the change from 1997–1999 to 2004–2006, the period marking the productivity slowdown, since the 2011–2013 period likely continues to reflect effects of the Great Recession. Our microdata-based approach closely matches BLS numbers for the 1997–1999 period, and we report an even larger deceleration in 2004–2006. We interpret the difference between 1997–1999 and 2004–2006 as reflecting the productivity slowdown and focus on decomposing it into changes in within-firm growth and changes in allocative components of aggregate growth.

II. Decomposing Productivity

There is a large literature on methods for decomposing aggregate productivity. We focus primarily on the Dynamic Olley Pakes method (hereafter DOP) of Melitz and Polanec (2015). Olley and Pakes (1996) showed that aggregate productivity can be decomposed into the unweighted average of firm-level productivity and a term that is proportional to the covariance between firm size and firm productivity (where we suppress time subscripts for convenience):

$$(1) \qquad P_i = \overline{p}_i + \mathrm{cov}(\theta_f, p_f),$$

where P_i is industry aggregate productivity, $\overline{p}_i$ is the unweighted average of (log) firm-level productivity for firms in industry i, θ_f is the share of industry employment accounted for by firm f, and p_f is the (log) labor productivity of firm f. The covariance term has been interpreted as a measure of allocative efficiency, or the degree to which higher-productivity firms have access to more resources. While this interpretation is more natural with a TFP productivity measure, Bartelsman, Haltiwanger, and Scarpetta (2013) show both theoretically and empirically that the Olley-Pakes decomposition applied to labor productivity yields similar inferences. They note that this inference is model dependent, but we adopt this interpretation in this short paper.

Melitz and Polanec (2015) extends the Olley Pakes method to include entry and exit in a way that allows for careful tracking of within-firm changes:

$$(2) \qquad \Delta P_i = \Delta \overline{p}_{i,C} + \Delta \mathrm{cov}_C(\theta_f, p_f) + \theta_{E2}(P_{E2} - P_{C2}) + \theta_{X1}(P_{C1} - P_{X1}),$$

where Δ indicates year-over-year log difference, C denotes continuer firms (those with employment in both years), $E2$ denotes entrants in the second year of the calculation, $X1$ denotes firms that exit after the first year, and $C1$ and $C2$ denote continuers in the first and second years, respectively. The first term in the expression, $\Delta \overline{p}_{i,C}$, represents average within-firm productivity growth for continuing firms; the second term, $\Delta \mathrm{cov}_C(\theta_f, p_f)$, represents the change in allocative efficiency among continuing firms; and the remaining terms represent the aggregate contribution of net entry. We calculate equation

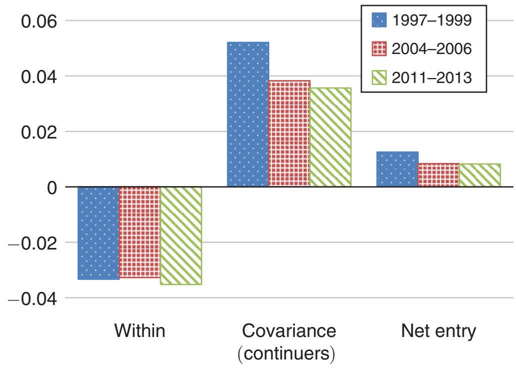

FIGURE 1

Source: Author calculations from RE-LBD.

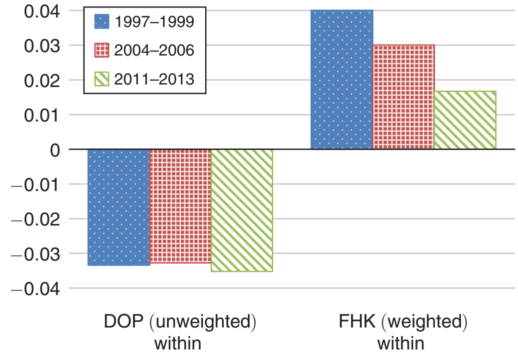

FIGURE 2

Source: Author calculations from RE-LBD.

(2) for each industry in each year and aggregate the annual components to the economywide level using fixed industry shares as described above. Figure 1 reports the resulting components of aggregate productivity growth.

On Figure 1, the first set of bars reports the average annual change in productivity within continuing firms for each noted time period, the second set reports the change in allocative efficiency among continuing firms, and the third set reports the contribution of net entry (see equation (2)). Notably, the within-firm contribution is consistently below zero; surviving firms see negative productivity growth on average. This negative contribution is roughly constant over time, suggesting that the productivity slowdown was not driven by reduced within-firm productivity growth on average. The covariance terms, reflecting the aggregate improvement in allocative efficiency among continuing firms, consistently account for the bulk of aggregate productivity growth; a step down is apparent between 1997–1999 and 2004–2006. Likewise, net entry makes a positive contribution to growth but steps down between 1997–1999 and 2004–2006 consistent with Alon et al. (2017); these authors show that declining entry has a significant cumulative effect on productivity over the 2000s. Strikingly, from an Olley-Pakes perspective the productivity slowdown between the late 1990s and the mid-2000s is accounted for by decelerating allocative efficiency, primarily among continuing firms but also in terms of net entry, rather than slowing improvements within firms.

It is notable that within-firm productivity growth is negative on average, but recall that this term is an unweighted average in the DOP framework. About 90 percent of firms have fewer than 20 workers, so the unweighted "within" term largely reflects the contribution of very small firms. These small firms account for only about 10 percent of total employment, so it is instructive to also examine weighted within-firm productivity growth. Indeed, in some dynamic decompositions from the literature (e.g., Foster, Haltiwanger, and Krizan 2001—henceforth, FHK) it is typical to compute a "within" term that is the weighted average of firm-level productivity growth among continuers, $\sum_f \theta_{f1} \Delta p_f$, where the weights θ_{f1} are equal to each firm's initial employment as a share of initial aggregate employment. The difference between this weighted approach and the unweighted DOP "within" term is given by

$$(3)\quad \sum_f \theta_{f1} \Delta p_f - \Delta \bar{p} = \sum_f \left(\theta_{f1} - 1/N\right) \Delta p_f,$$

where N is the number of firms. This difference will be positive only if within-firm productivity growth and the initial employment share (size) of the firm covary positively.

On Figure 2 we again report the "within" term from our DOP exercise and, in addition, we report the FHK weighted within-firm method. The FHK "within" term is always positive, but the unweighted (DOP) "within" term is negative and consistently less than its weighted counterpart (FHK); so larger firms must have higher within-firm productivity growth on average.

Rather than reflecting pure within-firm effects, then, the weighted (FHK) "within" term arguably also draws on allocative efficiency mechanics that, in the DOP framework, are instead counted in the changes in the covariance term. A comparison of the DOP and FHK terms on Figure 2 therefore reveals that (i) firms with higher productivity growth tend to be larger on average; (ii) this positive correlation has fallen over time and, in an FHK-type framework, accounts for some portion of the productivity slowdown; and (iii) differences in the changes of unweighted and weighted averages involve the interaction of allocative efficiency mechanics and within-firm productivity changes.

This last point suggests that it is wrong to conclude from Figure 1 that changes in within-firm productivity growth play no role in the productivity slowdown. While the unweighted "within" mean exhibits little change, underlying this unweighted mean is wide dispersion in within-firm productivity growth rates: we find that the interdecile range of within-firm productivity growth rates is about 100 log points. Based on equation (3), this dispersion combines with positive correlation between within-firm productivity growth and initial size to drive the positive difference between the weighted and unweighted means of within-firm productivity growth illustrated by Figure 2. To shed further light on this dispersion, Figure A2 of the online Appendix reports (unweighted) within-industry ninetieth percentile productivity growth rates, averaged across industries (using our usual time-invariant employment weights), by firm size class. The ninetieth percentile of within-firm productivity growth rates is high but declining for all firms and for every size class. Notably, the largest declines are seen among the largest size classes. In Figure A3 of the online Appendix, we conduct the same exercise using employment-weighted ninetieth percentiles and find that the weighted percentiles are only slightly smaller in magnitude and exhibit the same pattern of declines over time.

Taken together with Figure 2, Figures A2 and A3 imply that the productivity slowdown is partly driven by declines in the upper tail of the within-firm productivity growth distribution. Interestingly, there is a decline in the upper tail of the productivity growth distribution in every size class and in both unweighted and weighted terms. On figures A4–A6 of the online Appendix we also find that all size classes exhibit a large, positive difference between the weighted and unweighted "within" terms, and in all cases this gap declines from 1997–1999 to 2004–2006. These patterns manifest themselves partly in allocative efficiency terms since they both reduce the correlation between size and productivity growth and limit opportunities for further productivity-enhancing reallocation.

III. Conclusion

The evidence presented here advances the literature in three ways. First, decompositions of aggregate labor productivity growth suggest that impaired growth in allocative efficiency can account for the bulk of the productivity slowdown from the late 1990s to the mid-2000s. Current debates about the productivity slowdown focus on whether it reflects slowing technological improvement or increasingly imperfect measurement, but allocative efficiency is crucial for transmitting advances in technology and management practices into aggregate productivity growth. Decelerating allocative efficiency can constrain productivity growth even in the midst of rapid technological progress; alternatively, changes in technology may be influencing the pace of reallocation and possibly allocative efficiency measures. We also observe wide variation in within-firm productivity growth and growth slowdowns by firm size. This evidence should inspire a reevaluation of the productivity slowdown debate.[1]

Second, there are complex interactions between within-firm productivity growth and measures of allocative efficiency. The covariance between within-firm productivity growth and initial size has weakened, and the ninetieth percentile of the within-firm productivity growth distribution has fallen. The decline in the latter is more substantial for the largest firms, so in this respect it is difficult to draw clear distinctions between allocative efficiency and technological stagnation mechanisms.

Third, the evidence is consistent with the notion that post-2000 declining business dynamism has not been benign for American living standards but, instead, is closely related to

[1] See Andrews, Criscuolo, and Gal (2015) for another firm dynamics approach to the productivity slowdown.

slowing productivity growth. These results complement Decker et al. (2017), which finds that declining reallocation reflects a decline in the responsiveness of individual businesses to their productivity.

While the discovery of strong causal factors behind these patterns has thus far proven elusive in this literature, the accumulating evidence has narrowed the possibilities considerably while emphasizing the importance of the topic.

REFERENCES

Alon, Titan M., David Berger, Rob Dent, and Benjamin Pugsley. 2017. "Older and Slow: The Startup Deficit's Lasting Effects on Aggregate Productivity Growth." https://sites.google.com/site/davidwberger/research (accessed February 15, 2017).

Andrews, Dan, Chiara Criscuolo, and Peter N. Gal. 2015. *Frontier Firms, Technology Diffusion and Public Policy: Micro Evidence from OECD Countries.* Paris: Organisation for Economic Co-operation and Development (OECD).

Bartelsman, Eric, John Haltiwanger, and Stefano Scarpetta. 2013. "Cross-Country Differences in Productivity: The Role of Allocation and Selection." *American Economic Review* 103 (1): 305–34.

Byrne, David M., John G. Fernald, and Marshall B. Reinsdorf. 2016. "Does the United States Have a Productivity Slowdown or a Measurement Problem?" *Brookings Papers on Economic Activity*: 109–57.

Decker, Ryan A., John Haltiwanger, Ron S. Jarmin, and Javier Miranda. 2016. "Where Has All the Skewness Gone? The Decline in High-Growth (Young) Firms in the U.S." *European Economic Review* 86: 4–23.

Decker, Ryan A., John Haltiwanger, Ron S. Jarmin, and Javier Miranda. 2017. "Changing Business Dynamism and Productivity: Shocks versus Responsiveness." https://www.dropbox.com/s/exw6mnvgqufndrg/DHJM_AERPP_2017.pdf (accessed February 15, 2017).

Foster, Lucia, John C. Haltiwanger, and C. J. Krizan. 2001. "Aggregate Productivity Growth: Lessons from Microeconomic Evidence." In *New Developments in Productivity Analysis*, edited by Charles R. Hulten, Edwin R. Dean, and Michael J. Harper, 303–72. Chicago: University of Chicago Press.

Gordon, Robert. 2016. *The Rise and Fall of American Growth.* Princeton: Princeton University Press.

Haltiwanger, John, Ron S. Jarmin, Robert Kulick, and Javier Miranda. Forthcoming. "High Growth Young Firms: Contribution to Job, Output, and Productivity Growth." In *Measuring Entrepreneurial Businesses: Current Knowledge and Challenges*, edited by John Haltiwanger, Erik Hurst, Javier Miranda, and Antoinette Schoar. Chicago: University of Chicago Press.

Melitz, Marc J., and Saso Polanec. 2015. "Dynamic Olley-Pakes Productivity Decomposition with Entry and Exit." *RAND Journal of Economics* 46 (2): 362–75.

Olley, G. Steven, and Ariel Pakes. 1996. "The Dynamics of Productivity in the Telecommunications Equipment Industry." *Econometrica* 64 (6): 1263–97.

Syverson, Chad. 2016. "Challenges to Mismeasurement Explanations for the U.S. Productivity Slowdown." National Bureau of Economic Research Working Paper 21974.